

# CSC 212: Data Structures and Abstractions

## Balanced trees (part 1)

Prof. Marco Alvarez

Department of Computer Science and Statistics  
University of Rhode Island

Fall 2025



## Balanced search trees

- **Balanced search trees** are a type of trees that maintain structural invariants ensuring height  $h = O(\log n)$  for  $n$  nodes

- ✓ among the most useful data structures in computer science
- ✓ widely implemented in standard libraries:
  - Java: `TreeSet` and `TreeMap`,
  - C++: `std::set` and `std::map`
  - Python: no built-in, but available in libraries

- Examples of balanced trees:

- ✓ AVL trees, **Red-Black trees**, B-trees, Treaps, etc.

## Practice


- Assume a dictionary has  $n$  keys, and a book has  $m$  words
  - ✓ What is the time complexity of identifying which words from the book do NOT appear in the dictionary?
    - dictionary is represented as a BST and assume that  $h = O(\log n)$
    - book is represented as an array (vector) of strings, where each string is a word

2

## Red-black trees


- Red-black trees are BSTs that maintain a balanced structure by enforcing the following properties on the nodes:

- ✓ each node is colored either **red** or **black**
- ✓ the root node is always **black**
- ✓ **red** nodes cannot have **red** children (no two red nodes can be adjacent)
- ✓ **null nodes** are considered **black**
- ✓ every *root-to-null* path must have the same number of **black** nodes




4

## Examples



Red-black tree with implicit NIL leaves




Red-black tree with explicit NIL leaves

<https://www.happycoders.eu/algorithms/red-black-tree-java/>

5

## Practice

- Are these valid red-black trees? — (null nodes not shown)



6

## Analysis

- A red-black tree on  $n$  nodes has height  $h = O(\log n)$ 
  - after performing an insertion or deletion, the tree may temporarily violate red-black tree properties
  - to restore these properties, we efficiently modify the tree through:
    - rotations
    - recoloring nodes
- Equivalence to [2-3-4 Trees](#)
  - red-black trees correspond to 2-3-4 trees ([B-trees of order 4](#))
  - this correspondence provides intuition for understanding rebalancing operations and complexity analysis

7

## B-Trees (interlude)


## Multi-way search trees

- A multi-way search tree is a generalization of a BST where:

- ✓ each node can contain multiple keys (not just one)
- ✓ each node can have more than two children

### Properties

- ✓ keys within each node are **sorted** in increasing order
- ✓ the keys in the left subtree of a key  $k$  are less than  $k$ , and the keys in the right subtree are greater than  $k$



note that null pointers are illustrated as external nodes

Image credit: Data Structures and Algorithms in C++ 2e

9

## Balanced multi-way search trees

### Balanced multi-way search tree

- ✓ cap the number of children to a fixed number and **keep the leaf nodes at the same depth**
- ✓ the tree is **always balanced**
  - all leave nodes have the same depth
  - search, insertion, and deletion can be performed in  $O(\log n)$  time

### B-tree: specific type of a balanced multi-way search tree

- ✓ on a B-tree of **order  $m$** , each node, except the root, must have between  $\lceil b/2 \rceil$  and  $b$  children
- note there are differences in terminology (multiple “order” definitions)
- ✓ heavily used in databases and file systems to store large amounts of data (common orders: 1024, 2048, 4096, ...)

11

## Search on a multi-way search tree

- Perform **search** for 12, 17, 24, and 50 on the following tree
  - ✓ note that null pointers are illustrated as external nodes




Image credit: Data Structures and Algorithms in C++ 2e



10

Assume  $d$  denotes the maximum number of keys of any node of  $T$ , and  $h$  denotes the height of  $T$ . What is the cost of search?

## 2-3-4 tree

### A 2-3-4 tree (a.k.a. 2-4 tree) is a B-tree of order 4

- ✓ each node can have 2, 3, or 4 children
- ✓ i.e. all nodes must have at least 1 key and at most 3 keys, except the root node that can have 0 keys when the tree is empty



12

## Insertion (2-3-4 tree)


### Steps

- ✓ start at the root and traverse down the tree to find the appropriate leaf node
- ✓ if the leaf node has less than 3 keys, insert the new key in sorted order
- ✓ if the leaf node has 3 keys, split it into two nodes and promote the middle key to the parent node
  - insert the new key in the appropriate child node
  - if the parent node also has 3 keys, repeat the splitting process up to the root

### Tree remains balanced after each insertion

- ✓ all leaf nodes are at the same level

Insert 10, 20, 30, 40, 50, 60, 70, 80, 90, 100



13

## Practice

### Insert the following sequence into a 2-3-4 tree

- ✓ 15, 10, 25, 5, 1, 30, 45, 60, 100, 70, 80, 40, 35, 90

14

## Practice

### What is the max h of a 2-3-4 tree with n nodes?

- ✓ to maximize the height, we want to minimize the number of keys per node (instance of a worst-case)
- ✓ draw an example tree and express h in terms of n

15

## Practice

### What is the cost of search and insert on a 2-3-4 tree?

- ✓ worst-case scenario

16

## So far ...

- The cost of operations in a B-tree of order  $b$  is  $O(b \log_b n)$ 
  - ✓ insert, search, remove
  - ✓ small values of  $b$  make this cost optimal
- In practice ...
  - ✓ B-trees are widely used in databases and file systems to manage large amounts of data efficiently
  - ✓ useful for systems that read and write large blocks of data
    - B-trees can minimize the number of disk accesses required (much larger order values)

17


## Red-black trees

### Red-black trees $\Leftrightarrow$ 2-3-4 trees


- A 2-node in a 2-3-4 tree corresponds to a **black** node in a red-black tree



- A 3-node corresponds to a **black** node with one **red** child



- A 4-node corresponds to a **black** node with two **red** children




19

### Red-black trees $\Leftrightarrow$ 2-3-4 trees


- Red-black trees are **isometric** to 2-3-4 trees

- ✓ the number of black nodes on any *root-to-null* path corresponds to the number of levels of the 2-3-4 tree





- Every red-black tree can be transformed into an equivalent 2-3-4 tree and vice versa

- ✓ the relationship between the trees is not bijective (1-1)



20



## Example



- ✓ each node is colored either **red** or **black**
- ✓ the root node is always **black**
- ✓ **red** nodes cannot have **red** children (no two red nodes can be adjacent)
- ✓ null nodes are considered **black**
- ✓ every root-to-null path must have the same number of **black** nodes

21

## Example




- ✓ each node is colored either **red** or **black**
- ✓ the root node is always **black**
- ✓ **red** nodes cannot have **red** children (no two red nodes can be adjacent)
- ✓ null nodes are considered **black**
- ✓ every root-to-null path must have the same number of **black** nodes

22

## Practice

- Draw the red-black tree that corresponds to the following 2-3-4 tree



23