

CSC 212: Data Structures and Abstractions

Hash Tables (part 1)

Prof. Marco Alvarez

Department of Computer Science and Statistics
University of Rhode Island

Fall 2025



Can we do better?

3

Data Structure	Worst-case			Average-case			Ordered?
	insert at	delete	search	insert at	delete	search	
sequential (unordered)	$O(n)$	$O(n)$	$O(n)$	$O(n)$	$O(n)$	$O(n)$	No
sequential (ordered) binary search	$O(n)$	$O(n)$	$O(\log n)$	$O(n)$	$O(n)$	$O(\log n)$	Yes
BST	$O(n)$	$O(n)$	$O(n)$	$O(\log n)$	$O(\log n)$	$O(\log n)$	Yes
2-3-4	$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(\log n)$	Yes
Red-Black	$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(\log n)$	Yes

2

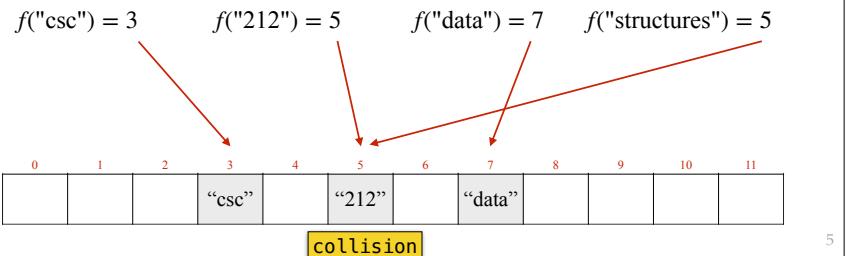
Random access memory

- **Random Access Memory (RAM)**
 - ✓ fundamental principle in computer science
 - ✓ enables constant-time access to any memory location, given its address — foundation for efficient array operations
 - ✓ arrays exploit RAM by mapping indices directly to memory addresses
- **Arrays in C++**
 - ✓ contiguous memory allocation
 - ✓ homogeneous elements (same data type)
 - ✓ fixed size (determined at creation, cannot be modified)
 - ✓ zero-based indexing
 - ✓ $O(1)$ random access

4

Hash tables

- A hash table is a data structure that implements an associative array
 - ✓ stores keys (**set**), or key-value pairs (**map**)
 - ✓ uses a **hash** function to compute an array index from a key
 - ✓ provides average-case $O(1)$ for search, insertion, deletion



Hash function

Hash function

- A hash function maps an input key to an integer value
 - ✓ hash value is then mapped to a valid array index using modulo
- Essential properties
 - ✓ **deterministic**: same key must always produce the same hash value
 - ✓ **uniform distribution**: hash values should be spread evenly
 - ✓ **efficient to compute**: execute in $O(1)$ time



Image credit: CS106B @ Stanford

7

Hash functions

- Space efficiency
 - ✓ supporting all possible keys directly would require enormous arrays => use a hash function to map keys to a much smaller array of size m (the **capacity** of the array)

```
// if hash() returns non-negatives
index = hash(key) % capacity
```

```
// if hash() returns any integer
index = abs(hash(key)) % capacity
```

The **load factor** α in a hash table is the ratio of N , the number elements, to M , the total capacity

8

Practice

- Which of the following tables is a better choice?

- what is the load factor α ?

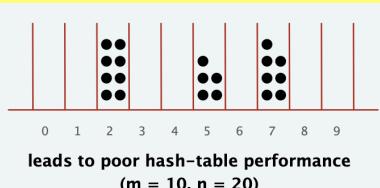
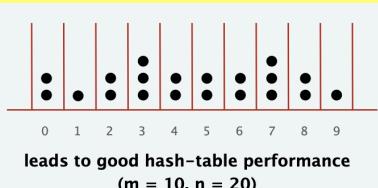


Image credit: COS 226 @ Princeton

9

Designing hash functions

- Importance of a hash function

- determines hash table's storage capacity
- directly impacts collision probability
- influences overall data structure performance

- Size selection strategies — mapping hashes into $[0, M - 1]$

- M is prime:** safe default, helps distributing keys more uniformly, minimizing collisions
 - hash $\% M$
- M is a power of two:** faster, enables fast modulo operation via bitwise AND
 - hash $\& (M - 1)$

11

Designing hash functions

- Design principles

- good hash functions satisfy: determinism, efficiency, and uniform distribution

- Hash functions on different data types

- integers:** may use the integer value as the hash value, or may apply transformations to break patterns
- floats:** convert to binary and treat as an integer, or manipulate the bits (e.g. XOR the mantissa and exponent)
- strings:** use the polynomial rolling hash ($31x + y$ rule) or other variants
- compound objects:** combine hash values of individual fields, may use the $31x + y$ rule or others

10

Collisions

- Definition

- occurs when two distinct keys hash to the same index in the table
 - inevitable consequence when $\alpha > 1$

- Resolution strategies:

- separate chaining:** each table slot contains a linked list (or other secondary structure) of all elements hashing to that index
 - insertion $O(1)$, search/delete $O(1 + \alpha)$
- open addressing:** all elements stored directly in the table, collisions trigger probe sequences
 - open addressing is more space-efficient than chaining, but it can be slower

12

Hash functions (beyond hash tables)

• Storing passwords

- ✓ never store plaintext passwords — use cryptographic hash functions

• File verification

- ✓ checksum mechanism — detect file corruption and provide data integrity validation
- ✓ verification process: generate and publish file hash, client recomputes hash of downloaded file, compare computed and published hashes

• Cryptographic hash functions

- ✓ one-way property: computationally infeasible to reverse the hash, find input producing specific hash, generate collision
- ✓ MD5, SHA-1, SHA-256, SHA-512, SHA3-512, ...

13

Separate chaining

Separate chaining

• Collision resolution mechanism

- ✓ utilizes a [linked list](#) at each hash table index (array of linked lists)
- ✓ guarantees $O(1)$ average-case and $O(n)$ worst-case search times
- ✓ supports dynamic memory allocation
- ✓ maintains key uniqueness constraint

• Core operations (assume a hash function h)

- ✓ **insert**: add a new key (or key/value pair) to the linked list at $h(key)$
 - insert at front for faster operations, no need to keep the keys on each list in sorted order
- ✓ **search**: search the linked list at $h(key)$
- ✓ **delete**: remove the key (or key/value pair) from linked list at $h(key)$

• Alternative collision resolution

- ✓ replace linked list with a self-balancing tree (e.g. Red-Black, AVL), guarantees $O(\log n)$ worst-case search

15

Practice

• Perform the following operations

- ✓ insert(L, 11), delete(D), insert(M, 12), delete(E), search(C)
- assume: insertions occur at front of the lists, hash(L)=3, hash(M)=0

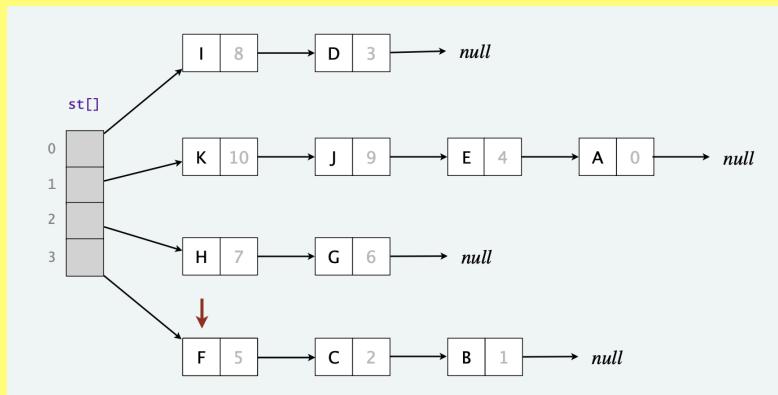


Image credit: COS 226 @ Princeton

16

```

#include <string>
#include <iostream>
#include <list>
#include <vector>

const size_t tableSize = 101;
typedef std::string Key;
typedef int Value;
typedef std::pair<std::string, int> HashEntry;
typedef std::list<HashEntry> HashChain;
typedef std::vector<HashChain> HashTable;

size_t hash(const Key& k) {
    size_t hashValue = 0;
    for (char ch : k) {
        hashValue = hashValue * 31 + ch;
    }
    return hashValue;
}

```

17

```

void insert(HashTable& table, const Key& key, const Value& value) {
    size_t index = hash(key) % tableSize;
    table[index].push_front({key, value});
}

Value* search(HashTable& table, const Key& key) {
    size_t index = hash(key) % tableSize;
    for (auto& entry : table[index]) {
        if (entry.first == key)
            return &entry.second;
    }
    return nullptr;
}

bool remove(HashTable& table, const Key& key) {
    size_t index = hash(key) % tableSize;
    auto& list = table[index];
    for (auto it = list.begin(); it != list.end(); ++it) {
        if (it->first == key) {
            list.erase(it);
            return true;
        }
    }
    return false;
}

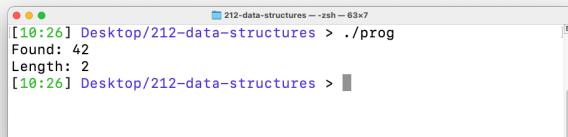
```

18

```

int main() {
    HashTable table(tableSize);
    insert(table, "example", 42);
    insert(table, "test", 84);
    insert(table, "csc", 126);
    insert(table, "212", 168);
    Value* val = search(table, "example");
    if (val) {
        std::cout << "Found: " << *val << std::endl;
    }
    remove(table, "example");
    remove(table, "test");
    size_t count = 0;
    for (const auto& chain : table) {
        count += chain.size();
    }
    std::cout << "Length: " << count << std::endl;
    return 0;
}

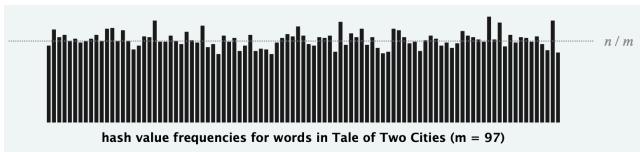
```



19

Analysis

- Uniform hashing assumption
 - assume the hash function is a good one, and all keys are uniformly distributed



- Load factor (α)
 - the ratio of the number of keys (n) to the number of slots (m)
- Time complexity
 - insert
 - $O(c)$ for all cases, if inserting at front of the list
 - search and delete
 - average case is $O(c + \alpha)$, where c is the cost of the hash function
 - worst case is $O(c + n)$, all the keys hash to the same index

$$\alpha = \frac{n}{m}$$

20

Considerations

Choices for α

- ✓ too small, results in excessive table size with inefficient space utilization
- ✓ too large, results in insufficient table size, increasing collision probability and degrading lookup performance

Typical values

- ✓ between 0.5 and 1.0 often provide a reasonable balance of space efficiency and lookup performance
- ✓ higher load factors (>1.0) remain functional but with degraded performance
- ✓ for performance-critical applications, conduct benchmarks with representative data sets to determine the optimal load factor

21

Data Structure	Worst-case			Average-case			Ordered?
	insert at	delete	search	insert at	delete	search	
sequential (unordered)	$O(n)$	$O(n)$	$O(n)$	$O(n)$	$O(n)$	$O(n)$	No
sequential (ordered) binary search	$O(n)$	$O(n)$	$O(\log n)$	$O(n)$	$O(n)$	$O(\log n)$	Yes
BST	$O(n)$	$O(n)$	$O(n)$	$O(\log n)$	$O(\log n)$	$O(\log n)$	Yes
2-3-4	$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(\log n)$	Yes
Red-Black	$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(\log n)$	Yes
Hash table (separate chaining)	$O(1)$	$O(n)$	$O(n)$	$O(1)^*$	$O(1)^*$	$O(1)^*$	No

(*) assumes uniform hashing and appropriate load factor

22